

```
1 #include "../common/common.h"
2 #include <cuda_runtime.h>
3 #include <stdio.h>
4
5 /*
6  * simpleDivergence demonstrates divergent code on the GPU and its impact on
7  * performance and CUDA metrics.
8  */
9
10 __global__ void mathKernel1(float *c)
11 {
12     int tid = blockIdx.x * blockDim.x + threadIdx.x;
13     float ia, ib;
14     ia = ib = 0.0f;
15
16     if (tid % 2 == 0)
17     {
18         ia = 100.0f;
19     }
20     else
21     {
22         ib = 200.0f;
23     }
24
25     c[tid] = ia + ib;
26 }
27
28 __global__ void mathKernel2(float *c)
29 {
30     int tid = blockIdx.x * blockDim.x + threadIdx.x;
31     float ia, ib;
32     ia = ib = 0.0f;
33
34     if ((tid / warpSize) % 2 == 0)
35     {
```

```
36     ia = 100.0f;
37 }
38 else
39 {
40     ib = 200.0f;
41 }
42
43 c[tid] = ia + ib;
44 }
45
46 __global__ void mathKernel3(float *c)
47 {
48     int tid = blockIdx.x * blockDim.x + threadIdx.x;
49     float ia, ib;
50     ia = ib = 0.0f;
51
52     bool ipred = (tid % 2 == 0);
53
54     if (ipred)
55     {
56         ia = 100.0f;
57     }
58
59     if (!ipred)
60     {
61         ib = 200.0f;
62     }
63
64     c[tid] = ia + ib;
65 }
66
67 __global__ void mathKernel4(float *c)
68 {
69     int tid = blockIdx.x * blockDim.x + threadIdx.x;
70     float ia, ib;
```

```
71     ia = ib = 0.0f;
72
73     int itid = tid >> 5;
74
75     if (itid & 0x01 == 0)
76     {
77         ia = 100.0f;
78     }
79     else
80     {
81         ib = 200.0f;
82     }
83
84     c[tid] = ia + ib;
85 }
86
87 __global__ void warmingup(float *c)
88 {
89     int tid = blockIdx.x * blockDim.x + threadIdx.x;
90     float ia, ib;
91     ia = ib = 0.0f;
92
93     if ((tid / warpSize) % 2 == 0)
94     {
95         ia = 100.0f;
96     }
97     else
98     {
99         ib = 200.0f;
100     }
101
102     c[tid] = ia + ib;
103 }
104
105
```

```
106 int main(int argc, char **argv)
107 {
108     // set up device
109     int dev = 0;
110     cudaDeviceProp deviceProp;
111     CHECK(cudaGetDeviceProperties(&deviceProp, dev));
112     printf("%s using Device %d: %s\n", argv[0], dev, deviceProp.name);
113
114     // set up data size
115     int size = 64;
116     int blockSize = 64;
117
118     if(argc > 1) blockSize = atoi(argv[1]);
119
120     if(argc > 2) size = atoi(argv[2]);
121
122     printf("Data size %d ", size);
123
124     // set up execution configuration
125     dim3 block (blockSize, 1);
126     dim3 grid ((size + block.x - 1) / block.x, 1);
127     printf("Execution Configure (block %d grid %d)\n", block.x, grid.x);
128
129     // allocate gpu memory
130     float *d_C;
131     size_t nBytes = size * sizeof(float);
132     CHECK(cudaMalloc((float**)&d_C, nBytes));
133
134     // run a warmup kernel to remove overhead
135     size_t iStart, iElaps;
136     CHECK(cudaDeviceSynchronize());
137     iStart = seconds();
138     warmingup<<<grid, block>>>(d_C);
139     CHECK(cudaDeviceSynchronize());
140     iElaps = seconds() - iStart;
```

```
141     printf("warmup      <<< %4d %4d >>> elapsed %d sec \n", grid.x, block.x,
142            iElaps );
143     CHECK(cudaGetLastError());
144
145     // run kernel 1
146     iStart = seconds();
147     mathKernel1<<<grid, block>>>(d_C);
148     CHECK(cudaDeviceSynchronize());
149     iElaps = seconds() - iStart;
150     printf("mathKernel1 <<< %4d %4d >>> elapsed %d sec \n", grid.x, block.x,
151            iElaps );
152     CHECK(cudaGetLastError());
153
154     // run kernel 2
155     iStart = seconds();
156     mathKernel2<<<grid, block>>>(d_C);
157     CHECK(cudaDeviceSynchronize());
158     iElaps = seconds() - iStart;
159     printf("mathKernel2 <<< %4d %4d >>> elapsed %d sec \n", grid.x, block.x,
160            iElaps );
161     CHECK(cudaGetLastError());
162
163     // run kernel 3
164     iStart = seconds();
165     mathKernel3<<<grid, block>>>(d_C);
166     CHECK(cudaDeviceSynchronize());
167     iElaps = seconds() - iStart;
168     printf("mathKernel3 <<< %4d %4d >>> elapsed %d sec \n", grid.x, block.x,
169            iElaps);
170     CHECK(cudaGetLastError());
171
172     // run kernel 4
173     iStart = seconds();
174     mathKernel4<<<grid, block>>>(d_C);
175     CHECK(cudaDeviceSynchronize());
```

```
176     iElaps = seconds() - iStart;
177     printf("mathKernel4 <<< %4d %4d >>> elapsed %d sec \n", grid.x, block.x,
178           iElaps);
179     CHECK(cudaGetLastError());
180
181     // free gpu memory and reset device
182     CHECK(cudaFree(d_C));
183     CHECK(cudaDeviceReset());
184     return EXIT_SUCCESS;
185 }
186
```